

Design of an Autonomous Racing Vehicle

Oluwatoni Ogunmade, Jack Xu, Tsugumi Murata, Adrian Malanaran, Angela Gu, Brian Tran,
Andrew Jin, Waleed Ahmed, Brian Kibazohi, Jun-Ha Jung

Abstract—This technical report describes the University of Waterloo’s team approach to the International Autonomous Robot Racing Competition(IARRC). We developed a system capable of real-time autonomous driving utilizing off-the-shelf parts including three cameras, an Inertial Measurement Unit(IMU), a quadrature encoder and a Hokuyo planar Lidar. The vehicle is a modified 1:10th scale Traxxas RC car with electronics to support battery monitoring, power distribution, remote control and emergency stopping. A robust and modular aluminum superstructure was built as a retrofit system for the vehicle, to provide a reliable and accessible platform for developing autonomous cruising system. A cutting edge neural net based traffic light detection system was added to improve robustness. A new path planner was designed relying on a Random-exploring Random Tree(RRT) planner, a Model Predictive Controller(MPC) for longer range planning and a trajectory rollout algorithm for short range evasive maneuvers.

In this paper we will focus on innovations in our overall approach by highlighting our computer vision, system architecture, software development practices, path planning and controls.

I. INTRODUCTION

The International Autonomous Robot Racing Competition (IARRC) is an annual event in which fully autonomous vehicles are tested in a race. Emphasis is placed on both safety and speed. These two objectives

MOST contributors are from the Departments of Mathematics and Engineering at the University of Waterloo.

are somewhat in opposition, and the challenge of the competition is to design a vehicle that satisfies both. The vehicles compete in two major races; a drag race focusing on high-speed control, and a circuit race focusing on safe maneuvering of a complex environment.

This report will focus on the mechanical, electrical and predominantly the software design of the robot. Advances over the designs from previous years will be demonstrated, along with the well-performing components which were retained. The following report is divided into the sections of problem definitions, mechanical design, design, software design, and conclusions.

II. PROBLEM DEFINITION

The design of the robot stems from the list of engineering requirements and objectives identified in the rules and implied in the details of the competition. The requirements are absolutely mandatory features the vehicles needs to meet. The objectives are performance targets derived from the races which the vehicle needs to complete.

A. Race Requirements

- The vehicle must not exceed 75cm long, 55cm wide, 60cm tall.
- The vehicle must not interfere with other robots operation, such as their sensing and navigation.
- The vehicle must detect the traffic light change reliably in varying conditions.

- The vehicle must be able to safely stop within 1 meter of crossing the finish line.
- The vehicle must have a wired and wireless emergency stop button.

B. Objectives

- The vehicle should be able to achieve 10 m/s at drag race.
- The vehicle should be able to drive higher than average 3 m/s overall during the drag race.
- The vehicle should be able to detect other robots and avoid collisions.
- The vehicle should be able to detect obstacles and lines at 30 Hz
- The vehicle must be able to survive a collision and inversion at 10 m/s.
- All parts should be easily accessible; can be replaced in fewer than 10 minutes.

III. MECHANICAL DESIGN

This year, a new mechanical retrofit system has been designed and built based on previous years' feedback. This year's mechanical system is much simpler, robust, modular, and rigid to provide a much more stable and easier developing platform.

A. Body

As previous years, the body of the robot is designed to be a retrofit system for a Traxxas RC monster truck. To make this year's chassis much more robust and rigid to protect on-board hard-wares, while keeping its simplicity and flexibility. The entire chassis is mainly made of aluminum angle extrusions, sheet, and 3D printed ABS components to provide a secure and light-weight system. The main body (as shown in Figure 7) can be composed

with three major parts: a wire frame superstructure, a build plate, and other mounting components.

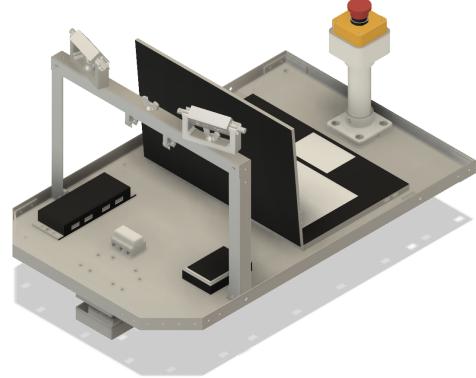


Fig. 1: The Fusion CAD model of the body

B. Wire Frame Superstructure

The wire frame superstructure is the core component that bridges and secures other components. The wire frame is mainly made of aluminum materials. It is composed with three parts: a retrofit mounting system, a protective cage, and a camera rig. For retrofit mounting system, it is built with two aluminum clevis with two angle extrusions attached horizontally. The protective cage is made of aluminum in a trapezoid shape. The base cage can provide far enough strength to prevent the build plate away from any torsion or shear stresses from external forces or impacts. Similarly, the camera rig is also made of aluminum angle extrusions to provide a stable and flat platform for cameras.

C. Build Plate

The main part of the build plate is a $\frac{1}{16}$ in. aluminum sheet, attached inside the protective cage securely. Since the build plate is entirely independent from the superstructure, it can be freely accessed and modified easily.

and independently without dissembling other structures. This can make the robot much easier to maintain and develop separately without complicated disassembly processes. On the bottom side of the plate, the lidar is also directly attached onto the sheet to remain protected and rigidly, while providing a clear detecting region (as shown in Figure 2). On the top side of the plate, hardware components such as a MCU (Micro Controller Unit) and a laptop can be safely mounted in place during the race.

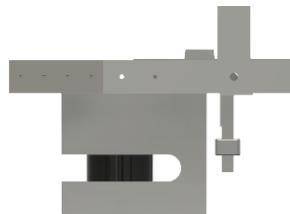


Fig. 2: Detailed view of the Lidar mount in the Fusion from the left side

triangle friction disk is designed to fix the camera while allowing the team able to freely adjust the camera angles to optimal positions. Camera mounts can then be placed onto the camera rig and fixed with bolts and nuts.

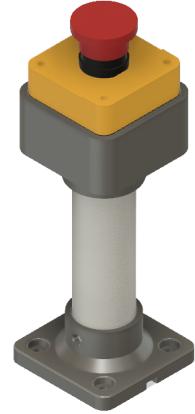


Fig. 3: E-Stop mount in Fusion



Fig. 4: Camera mount in Fusion

D. Other Mounting Components

Other mounting components are mainly about the mounting system for mechanical E-Stop, and cameras. To keep the E-Stop in the required height and accessible area, a pillar with the combination of 3D printed mounting base and housing and an aluminum pipe was developed to elevate the E-Stop at the right height (as shown in Figure 3). In addition, the E-Stop pillar is installed right onto the wire frame superstructure to keep it securely and rigidly in place. For cameras, web cameras used in the previous competitions are kept to be used this year. To reduce the wobbliness of camera due to weak mounting mechanisms and sudden motion, a camera housing and mounting systems (as shown in Figure 4) are designed in Fusion and 3D printed with ABS filaments. At the surface of each joint a rough

IV. ELECTRICAL DESIGN

ADD A SUMMARY OF THE ELECTRICAL DESIGN LAYOUT

A. Power

Two 12V 5800mAh Li-Po batteries, connected in series, provide the main 24V power rail to the robot. This 24V is directly used by the brushless motor. For other components, the power is regulated into different voltage rail, depending on how much the respective component

requires. The Hokuyo 2D planner lidar requires a 12V power, and therefore one regulator step downs from 24V into 12V. Additionally, the Arduino and the USB hub requires a 5V, because of that another regulator step downs from 12V into 5V. To monitor the status of battery, a battery management is integrated on the Arduino Shield. The BMS (battery management system) monitors the voltage of the battery to check if the current battery level and reports that to the main computer.

B. Arduino Shield

The arduino shield PCB is built from scratch using CAD tool to provide easy and simple integration of the peripherals with the Arduino Mega. The PCB consists of connection with the RC transmitter, servo motor, brushless motor, BMS, IMU, and the 8bit-AVR MCU which is dedicated for encoder counter. This encoder counter is the new design for the PCB, and the previously used encoder counter IC is replaced due to the expensive cost. The encoder counter MCU receives the signal from the encoder, and outputs the encoder count via I2C to the Arduino mega. The PCB also contains the pull up resistors required for the I2C communication, and hence the sensors requiring I2C interface such as IMU can be directly plugged in.

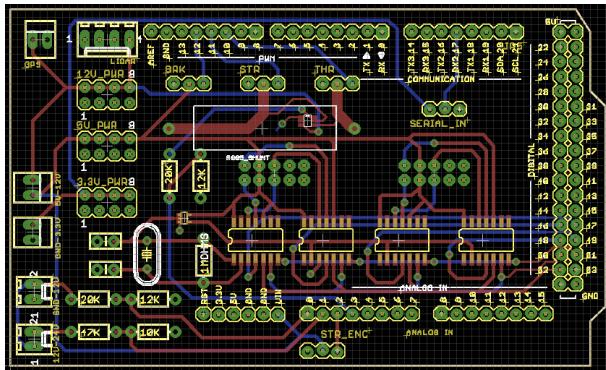


Fig. 5: The Circuit Layout of the Arduino shield

The Arduino shield allowed for the direct interfacing of motor encoders, counters and voltage regulators for the Arduino. This allowed the PID to be moved onto the Arduino, and a very high control loop frequency to be supported. Formerly, the PID controller was in a ROS node on the main PC.

C. Motor Controls

The motors used for powering the vehicle comprised two steering servos and one brushless drive motor. The drive motor was controlled by a Castle Creations Electronic Speed Control unit (ESC). These motors are all controlled by the Arduino via PWM servo commands. The Arduino utilizes the PID controller developed from the previous year to maintain accurate velocity of the drive motor, assuming no slip condition on the wheels. In the event the Arduino PWM signal is neutral or missing, the ESC causes the drive motor to brake when powered.

D. Sensors

The sensor suite on the vehicle comprises of both external and internal sensors. The external sensors include the 3 USB webcams and the Hokuyo 2D planar scanning Lidar. The internal sensors are the motor encoder, IMU and the battery monitor.

As described in the electrical design section, encoder counter is built from scratch using a 8-bit AVR MCU. The MCU which is programmed in low level C is capable to detect 512 encoder counts per revolution. The two encoder pulses are each detected as an external interrupt, and depending on the how the pulses change, the rotation direction and speed of the motor can be determined. The MCU communicates with the main Arduino via I2C, and returns the encoder count as unsigned 16 bit when the Arduino requests it. The calculation to convert from

the encoder count to distance and speed is done on the Arduino side. The distance and speed of the robot is then used during the motion control of the robot such as PID.

$$V_{car} = \frac{d}{dt} \left(\frac{Count_{encoder}}{256} \right) \cdot (R_{gear}) \cdot (\pi D_{wheel}) \quad (1)$$

The Hokuyo 2D lidar provides the vehicle with range data for any obstacles within 30m, in a 270° plane in front of the vehicle. The refresh rate for the full scan is 40Hz, with a practical accuracy of 5% with 0.5° resolution. This allows avoidance of any physical obstacle within detection range at high speed.

The vision system comprises 3 cameras, two for line detection and one for traffic light detection. They were selected because they were cheap and efficient in terms of computer and power resources. The occupancy grid is created from a combination of the lidar data and camera data, outlined in the software section.

E. Processing

The on board computer remains the same as it was last year. It remains a Lenovo Yoga 700 with a powerful 6th generation i7 chip from Intel, 8GB of RAM, and a discrete Nvidia graphics card. The RAM and graphics card allow more complex and parallel computations to be performed for algorithms which do the image processing and path planning. The PC communicates directly with all of the external sensors, and uses a serial connection with the Arduino to interface with the motors and internal sensors.

The low level state machine, PID motor control and system enable were all retained from the previous year. With the 3 states being the manual control mode, autonomous mode and E-stop mode. Switches on the remote controller are used to transition between states.

V. SOFTWARE DESIGN

The improvements over last year's software fall into three main categories: Architectural, Software development and testing, computer vision improvements and optimizations to sustaining higher speeds. These improvements allowed the team to radically improve our throughput and code quality while simultaneously increasing the performance of our robot.

A. Software Development and Testing

A holistic approach was taken to identify the limitations of the teams software development velocity. The main culprits were poor code documentation and lack of accessible testing. To improve our code documentation the code base was completely refactored with a focus on clearer comments and aligning with a coherent coding style. This greatly improved readability and allowed new team members to get up to speed very quickly. To ensure that the new standards were carried forward, a code review was made compulsory before all code submissions. This slowed development slightly but improve the quality of the code being submitted and ensured that only tested code was flowing into the competition code. To improve our software testing, there were two major improvements made over the previous year. Firstly a simulation of the race and our competition vehicle were made. This simulation was developed using gazebo and ROS. Three race tracks were modelled in the simulation, the drag race, circuit race and 60 meter long oval track. The tracks were simulated to closely resemble the race environment with traffic cones and lane markings.

The car was simulated with all the sensor currently on board the physical vehicle. The simulations allowed the software team to test both individual features and the

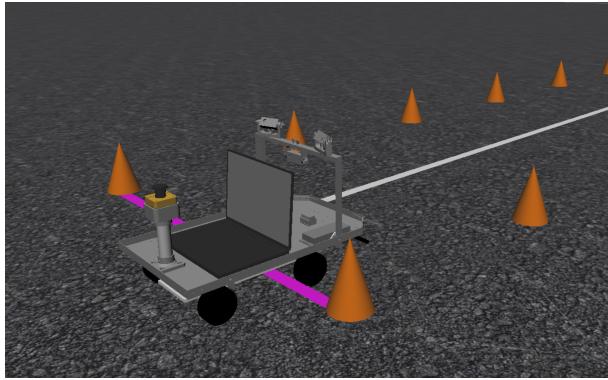


Fig. 6: Simulated Robot

entire system while the physical vehicle was still being built.

Secondly, we relied on recorded data from the past years competition to validate our code. A repository was created, containing the sensor data from previous competitions. This was used to validate our computer vision software and demonstrate it's superiority to software from the previous years.

B. Architectural

In the previous structure each node was responsible for keeping track of the current state of the robot with no node responsible of keeping track of metrics or overseeing the operation of the vehicle. The new software architecture is illustrated below, with the addition of the supervisor node at the helm of the system.

The sensing software consists of device drivers and ROS nodes that publish sensor data to other nodes to be consumed. Most of the software in this section is provided by vendors or the open source community. The path planning and controls sections are responsible for determining where to go and keeping the car on that path. The signal detection section consists of ROS nodes responsible for detecting important signals like a traffic light changes and the line signifying the end of the race.

They communicate these events to the supervisor which then decides on the course of action. The supervisor node is a module designed to manage high level decisions and keep track of useful race metrics. It provides an interface for signal detection nodes to indicate that an event has occurred. Using it's knowledge of the current state of the robot, it decides when to start and stop the car. To start and stop the car it coordinates with a twist multiplexer node to prioritize which twist messages get published to the robot. The twist multiplexer is a third party node that takes various twist message inputs, and based on the priority of the messages and any locks, will pick one to output to the ROS topic that the vehicle controller subscribes to. Currently, there are three sources of twist messages: joystick, path planner, and the supervisor node itself. The supervisor node's command topic is simply used to sending null vectors to stop the robot. Additionally, the supervisor node publishes a Boolean to two topics that the twist multiplexer looks for to set priority locks on the incoming twist messages. One topic is used to lock out all messages except for the supervisors topic, which is used to bring the robot to a complete stop. The other topic is to remove this lock, allowing all messages through, at which point the message with the highest priority will go through. The twist message priority is as follows: supervisor, joystick, path planner. The supervisor node also keeps track of some useful variables; namely the race type, start and end times, lap count, average speed, and battery life. When the race is finished, a few of these variables are stored and saved into a text file. These variables are used to make important decisions regarding when to start and stop the race. The race is started when a traffic light signal is received, and is ended when either the lap count is reached for that race (1 for drag, 3 for circuit), which

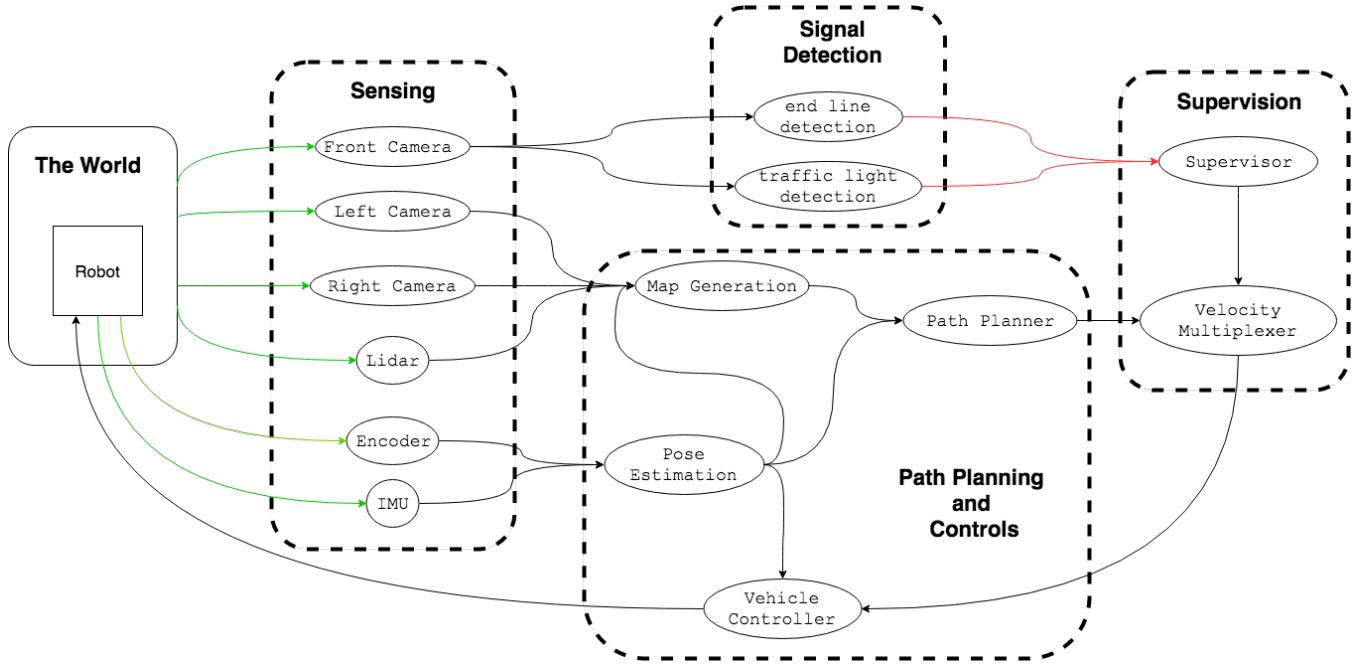


Fig. 7: The SolidWorks model of the body

is incremented by an end line detection signal, or if the battery life dips below 10

C. Computer Vision

This competition year huge strides we're made in improving the robustness and functionality of our computer vision software. To improve the robustness of our lane detection system a shadow removal algorithm was added as a pre-processing step on each of our cameras. To improve our traffic light detection in all conditions a machine learning classifier was added to help detect traffic lights in images. End line line detection was improved considerable over past years as well largely due to the addition of a new hysteresis filter on the camera stream.

1) Shadow Removal: In last years competition, the robot had difficulty in reliably tracking lane lines under variable lighting conditions, such as when buildings or robots cast shadows onto the track. As a result, this years

team implemented a shadow detection and removal pre-processing step, based on work done by Murali et al[cite], in order to address this issue. First, to detect shadows, the algorithm calculates the mean value of each plane in the Lab colour space and applies a threshold on the L and b values. The Lab colour space was designed to be uniform with respect to human vision, so it could better model visually perceived changes, such as variations in lightness that would indicate a shadow region. An example is shown in Figure X1.

To prepare for shadow removal, the shadow mask is dilated to smooth out holes and sharp edges and then separated into shadow regions. For each shadow region, the algorithm calculates the ratio of red, green, and blue values between the area inside the mask and the area just outside the mask. To remove shadows, the shadow regions RGB values are multiplied by these ratios. The results can be seen in Figure X2.

The last, and slowest, step is to correct for over-

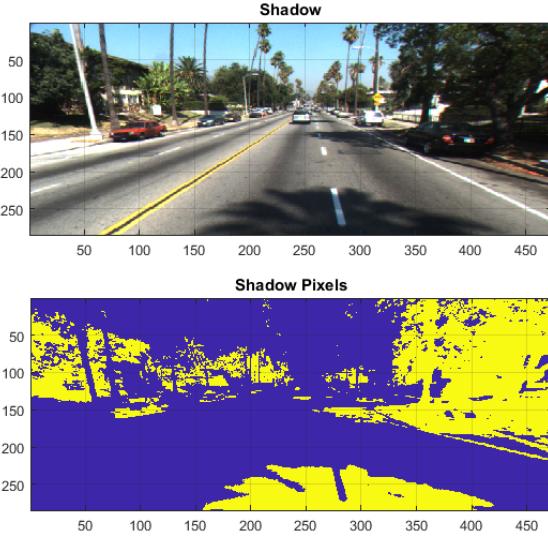


Fig. 8: Shadow Detection Using Lab Colour Space

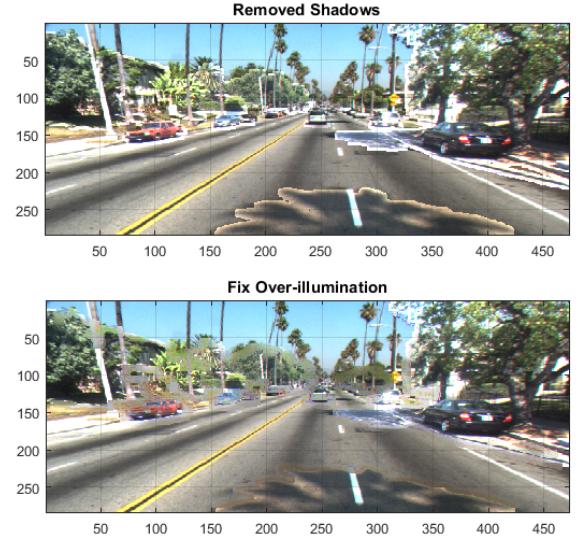


Fig. 10: Correction for Over-illuminated Shadow Region Edges

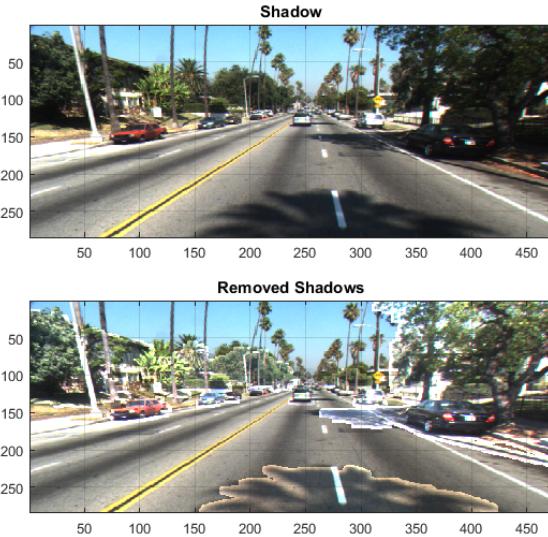


Fig. 9: Shadow Removal Using RGB Ratios

illuminated edges in the shadow regions, since the outer edges of shadows tend to be brighter than the inner region. To do this, a median filter is applied on the edges of each shadow region. Figure X3 shows this process.

There were two issues the team encountered with this shadow detection and removal approach. The first was that the over-illumination correction took the most time to run, but often blurred out lane lines through its

median filter, especially if the shadow was spotty and had edges in its interior. However, the team did not expect to encounter this situation in competition and the resulting images were often more useful for line tracking than the originals. An alternative method to correct for over-illumination could be considered in the future. Second, since the shadow detection relies on a mean value for the lightness (L) over the entire image, an frame with extreme highlights can cause the algorithm to misidentify shadows and over-illuminate entire regions. Again, this was not expected to be seen in competition since this issue was most common with a blown-out sky region and the robots cameras are directed towards the ground. These two issues are shown in Figure X4 and X5.

Figure X6 displays the final result, showing the lane detection before and after the shadow removal pre-processing step for a test image.

2) Traffic Light Detection: The traffic light detection process is initiated starting with availability of image data supplied by the camera sensors, which is relayed

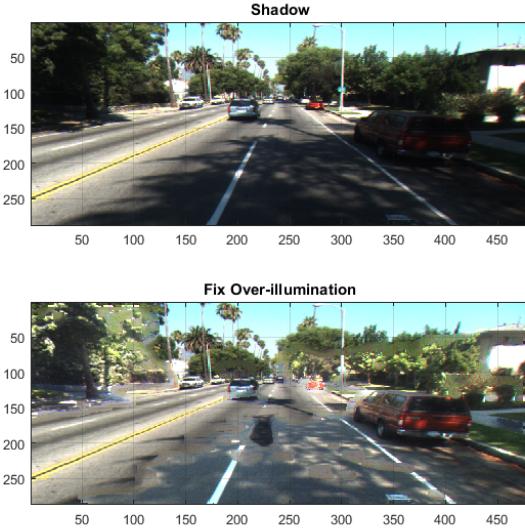


Fig. 11: Over-illumination Correction on Shadows with Inner Edges

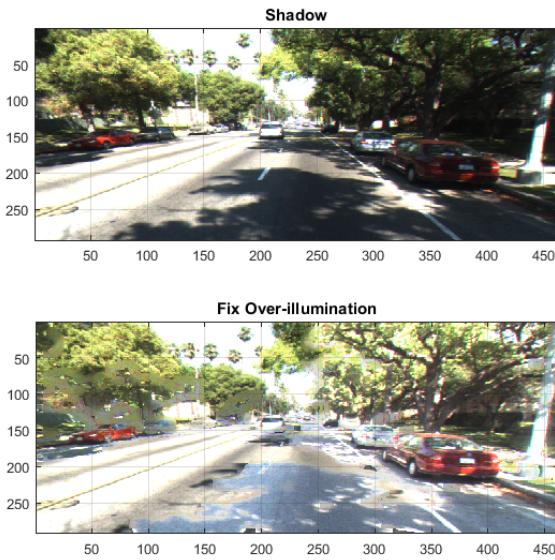


Fig. 12: Final Result of Shadow Removal on an Overexposed Image

in the form of a mutable message data structure to the system nodes. Detection of the traffic light object is performed using YOLO (You Only Look Once), a neural net based object classifier which is integrated as a separate system node responsible for identifying and

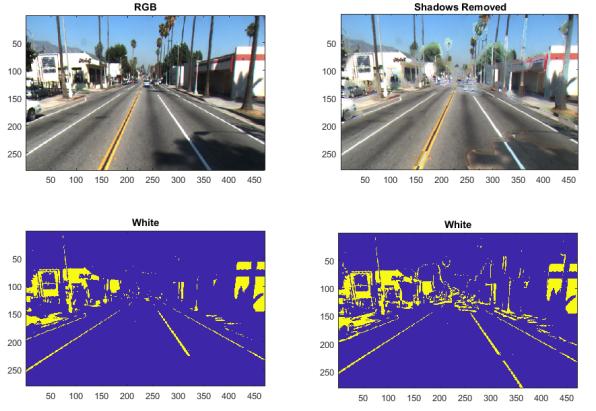


Fig. 13: White Line Detection Before and After Shadow Removal

outlining the location of specific objects within an given image [1]. As an example, Figure 1.0 shows YOLO detecting a traffic light object from camera image.

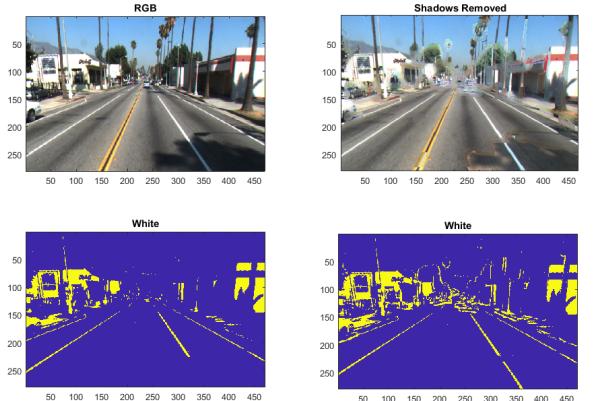


Fig. 14: White Line Detection Before and After Shadow Removal

Once a traffic light object is detected, the camera image is cropped to the bounding box of the traffic light, which is then sent over to the traffic light processor node. Within this section of the system, the image, existing in RGB color space is converted into a hue, value, saturation (HSV) space. Gaussian blur is also applied to the image to smoothen out irregularities and remove

noise to increase analysis accuracy. The state of the traffic light is detected by applying a filter that isolates pixels that correspond to a certain color range to distinguish red and green colors in a diverse set of lighting conditions. The traffic light color state is determined based off of the prominence in either color value. Upon detection of a green traffic light, a service call is invoked, signalling the start of the race. Testing is performed under multiple lighting conditions with various traffic light placements to determine best threshold parameters to ensure detection robustness.

3) End line Detection: At the end of each race there is a magenta line to indicate that. The algorithm used to detect the end line transforms camera images first from RGB to HSV. The threshold required to detect the end line is tuned to find a magenta hue, corresponding to a specific number range. Gaussian blur is used to reduce small detection errors. Finally, blob detection is used to filter for minimum area and non-circularity. This detection algorithm is inserted between a hysteresis counter procedure, that increments when detection is true and decrements when detection is false. This prevents false positive detection of the algorithm. To improve performance, the algorithm calls upon the supervisor node to decide upon the action to take, deciding based on the response to stop the entire robot after finishing the race.

D. Path Planning and Controls

The path planning framework was overhauled from the previous year. Whereas last year's path planner was entirely reactive, the new path planner incorporates a look ahead distance whenever lane markings are present. It uses these lane markings to pick a point ahead in space to plan a path to, this helps it sustain higher speeds

for longer stretches of time. The path is continually reassessed and planned and a fairly low frequency of 1 Hz. If any obstacles are detected along the path closer than a certain threshold the path planner falls back to the trajectory roll-out algorithm which can run at a much higher frequency of up to 20Hz. To remain on the path that was selected the robot utilizes an MPC. This is updated with the robots estimate of its current position by a sensor fusion ROS node. The MPC computes the error between where the robot is and where it should be, often referred to as the cross track error(CTE) and generates steering and throttle commands to follow the path.

1) Medium Range Planning: To plan a feasible path to a position farther ahead an RRT planner is employed. When provided a point along the track farther ahead of it. The planner takes the kinematic model into account and generates an obstacle free path for the robot to follow.

2) Close Range Obstacle Avoidance: The trajectory rollout algorithm from previous year is preserved and it is responsible for engaging in close range evasive maneuvers when unforeseen obstacles arise. These obstacles could be other cars on the track or humans that might cross the track. The algorithm was parallelized on a GPU to improve the performance over previous years.

3) Sensor Fusion and Controls: The velocity information from the encoder and orientation and acceleration information from IMU were fused together using an Extended Kalman Filter(EKF) to provide the path planner with the vehicles pose. The MPC uses the vehicle pose and the generated path to generate steering and throttle commands for the robot. The MPC is superior to our previous sole PID controller because it

VI. CONCLUSION

Thus far, the success of the current design can only be assessed relative to the performance of the previous vehicle. Many weak points were addressed, and significant performance improvements were made. The body was redesigned to be more durable, repairable and serviceable. The traffic light detection strategy was made much more robust, and has been tested successfully in varied environments. The path planner was improved to enable sustained high speed driving and retain vehicle responsiveness. The teams software development practices were updated and enforced. ADD FINAL SENTENCES.

ACKNOWLEDGMENT

The authors would like to thank members of the WAVE lab who provided hardware used on the robot. They would also like to thank Matthew Post for his mentorship and continued support of the team.

Additional note - All work done on changes to mechanical, electrical and software design was performed by team members who had no former experience with the competition, while returning members took on a mentoring role. The new members should be especially noted for their dedication, work ethic and ingenuity. They made this happen.

REFERENCES

- [1] J. Lee et al, *Design of Vision Assisted Navigation Ground Unmanned Advanced Racing Drone (VANGUARD)*, University of Waterloo, 2015.
- [2] J. Redmond, R. Girshick, A. Farhadi, and S. Divvala, *You Only Look Once: Unified, Real-Time Object Detection* Washington, tech, 2015
- [3]